

# ML Cheat Sheet

## Closed-form solutions

Closed-form solutions exist only for simple models with simple loss (typically quadratic) and assumptions. Closed-form solutions examples: Ordinary Least Squares(OLS), Ridge Regression, Weighted Least Squares, Gaussian Maximum Likelihood Estimation(MLE).

Matrix inversion costs  $O(n^3)$ , which does not scale. If  $X^T X$  is singular or ill-conditioned, the solution does not exist or is unstable.

## Linear Model

To find the minimum of the loss function, we compute partial derivatives:

$$\frac{\partial \ell(\beta_0, \beta_1)}{\partial \beta_1} = 0, \quad \frac{\partial \ell(\beta_0, \beta_1)}{\partial \beta_0} = 0.$$

Solving these equations gives:

$$\beta_1 = \frac{\sum x_i y_i - n \bar{x} \bar{y}}{\sum x_i^2 - n \bar{x}^2} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2} = \frac{\text{Cov}(x, y)}{\text{Var}(x)}.$$

The intercept is:

$$\beta_0 = \bar{y} - \beta_1 \bar{x}.$$

Consider a linear model with multiple outputs:

$$y_i = \sum_{j=1}^n x_{ij} w_j$$

For a fixed weight  $w_k$ :

$$\frac{\partial y_i}{\partial w_k} = \frac{\partial}{\partial w_k} \left( \sum_{j=1}^n x_{ij} w_j \right) = x_{ik}.$$

Stacking all derivatives together gives:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{w}} = X^\top.$$

For a function:

$$f : \mathbb{R}^n \rightarrow \mathbb{R}^m, \quad f(\mathbf{x}) = \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{bmatrix},$$

the derivative is the **Jacobian matrix**:

$$J = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{pmatrix}.$$

**Each row corresponds to one output, and each column corresponds to one input dimension.**

For a scalar-valued function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ , the **Hessian matrix** captures second-order derivatives:

$$H(f) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}.$$

The Hessian describes **curvature(曲率)**. In optimization:

- Positive definite(正定) Hessian  $\rightarrow$  unique global minimum
- Indefinite(不定) Hessian  $\rightarrow$  saddle points(鞍点) may exist

For squared loss in linear regression, the Hessian turns out to be constant(常数) and positive semidefinite(半正定), which guarantees convexity(凸性).

Rewrite the model in vectorized form:

$$\hat{\mathbf{Y}} = XW,$$

where the first column of  $X$  represent the bias (intercept) term. Using squared error, the loss becomes:

$$\ell(W) = \|Y - XW\|_F^2.$$

This **compact expression encodes all data points simultaneously**. Taking the derivative with respect to  $W$ :

$$\frac{\partial \ell(W)}{\partial W} = -2X^\top(Y - XW),$$

which means **the direction of steepest descent in parameter space**. Setting the gradient to zero yields the **normal equation**:

$$X^\top XW = X^\top Y,$$

which is **to characterize the optimal solution**.

If  $X^\top X$  is **invertible**, the optimal solution is:

$$W^* = (X^\top X)^{-1}X^\top Y.$$

This expression is known as the **least-squares solution**, and  $(X^\top X)^{-1}X^\top$  is the **pseudoinverse of  $X$** . Key interpretation:

- **The solution projects  $Y$  onto the column space of  $X$ .**
- **Learning is equivalent to geometric projection in feature space.**

Figure 1 below shows the three different datasets (A, B, C) and the same linear model used to predict each of these datasets.

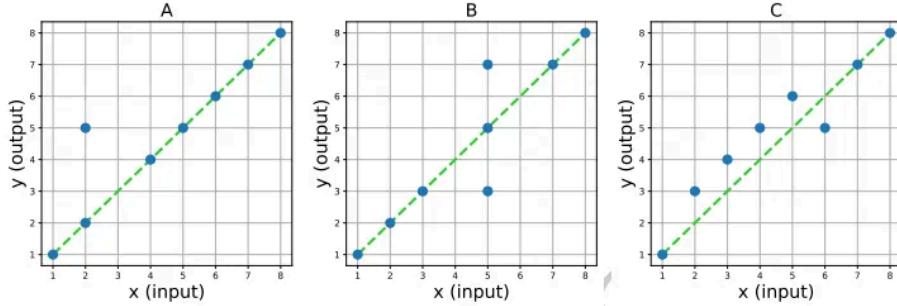


Figure 1: 3 Datasets (A, B, C), and a linear model represented by the dashed line (from  $\mathbb{R}^1$  to  $\mathbb{R}^1$ )

**Question 1** Choose the correct ordering with respect to MSE (Mean Squared Error) loss

- $MSE_B > MSE_A > MSE_C$
- $MSE_B > MSE_C > MSE_A$
- At least two of the MSE losses are equal.
- $MSE_A > MSE_C > MSE_B$
- $MSE_C > MSE_B > MSE_A$
- $MSE_A > MSE_B > MSE_C$
- $MSE_C > MSE_A > MSE_B$

**Solution:**  $MSE_A = 9/8$ ,  $MSE_B = 8/8$ ,  $MSE_C = 5/8$

**Question 2** Choose the correct ordering with respect to MAE (Mean Absolute Error) loss

- $MAE_A > MAE_C > MAE_B$
- $MAE_B > MAE_C > MAE_A$
- $MAE_A > MAE_B > MAE_C$
- $MAE_C > MAE_A > MAE_B$
- $MAE_B > MAE_A > MAE_C$
- At least two of the MAE losses are equal.
- $MAE_C > MAE_B > MAE_A$

**Solution:**  $MAE_A = 3/8$ ,  $MAE_B = 4/8$ ,  $MAE_C = 5/8$

## Likelihood and log-likelihood

For a single observation  $y$ , conditioned on  $x$  and  $w$ , the likelihood is

$$p(y | x, w) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2\sigma^2}(y - x^\top w)^2\right).$$

For i.i.d. samples, the joint likelihood is the product over all data points. After we immediately take logs:

$$\log p(y | X, w) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - x_i^\top w)^2.$$

The first term does not depend on  $w$ . Therefore:

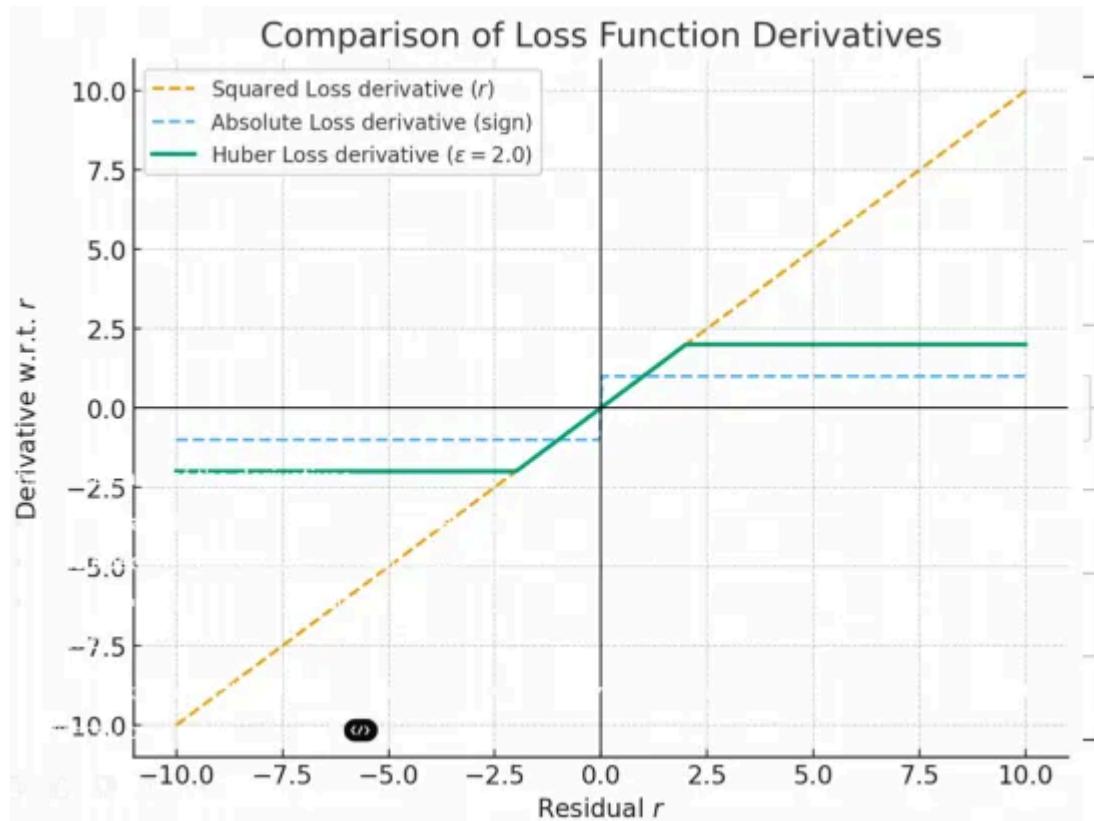
$$\arg \max_w \log p(y | X, w) \iff \arg \min_w \sum_{i=1}^n (y_i - x_i^\top w)^2.$$

Maximizing the log-likelihood under Gaussian noise is equivalent to minimizing the sum of squared errors.

Choosing a loss function is choosing a statistical assumption. Every loss corresponds to an implicit noise model.

- Gaussian noise  $\Rightarrow$  squared loss
- Laplace noise  $\Rightarrow$  absolute loss / L1 Loss
- Heavy-tailed noise  $\Rightarrow$  robust losses

## Huber loss: Bounded like L1 and smooth like L2



Define the residual

$$r_i = f_\beta(x_i) - y_i.$$

The Huber loss with parameter  $\varepsilon > 0$  is

$$\ell_\varepsilon(r) = \begin{cases} \frac{1}{2}r^2, & |r| \leq \varepsilon \\ \varepsilon \left(|r| - \frac{\varepsilon}{2}\right), & |r| > \varepsilon \end{cases}$$

$\varepsilon$  determines the transition point between L2 and L1 behavior.

First compute the derivative of the scalar Huber loss with respect to its argument  $r$ :

$$\ell'_\varepsilon(r) = \begin{cases} r, & |r| \leq \varepsilon \\ \varepsilon \operatorname{sign}(r), & |r| > \varepsilon \end{cases}$$

Now apply the chain rule:

$$\nabla_\beta r_i = x_i.$$

So the gradient of the full objective is

$$\nabla_\beta L(\beta) = \sum_{i=1}^n \ell'_\varepsilon(r_i)x_i = \sum_{i:|r_i|\leq\varepsilon} r_i x_i + \sum_{i:|r_i|>\varepsilon} \varepsilon \operatorname{sign}(r_i) x_i.$$

Why does the gradient of the Huber objective involve a sum of terms of the form  $\ell'_\varepsilon(r_i)x_i$ ? Because of the chain rule applied to  $r_i = x_i^\top \beta - y_i$

In robust statistics language, Huber loss implements **influence capping**.

There is **no closed-form solution** for Huber regression in general. The moment you leave pure L2, you also leave normal equations behind.

Because the objective is piecewise-defined and not purely quadratic.

## Multivariate Linear Regression

Instead of a single feature  $x$ , we now consider multiple features:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k.$$

Geometrically, with  $k$  features, the model is a hyperplane in  $\mathbb{R}^{k+1}$ .

For a dataset with  $n$  samples and  $k$  features, predictions can be written as:

$$\hat{y}_i = \omega_0 + \omega_1 x_{i1} + \cdots + \omega_k x_{ik}, \quad i = 1, \dots, n.$$

The regression problem is equivalent to solving the linear system:

$$XW = Y.$$

When the number of data points exceeds the number of parameters ( $n > k$ ), the system is **overdetermined**. In general:

- There is **no exact solution** satisfying  $XW = Y$ .
- The system is inconsistent due to noise.

The standard approach is **least squares**, minimizing:

$$\|XW - Y\|_2^2.$$

If  $X$  has full column rank, the solution is:

$$W = (X^\top X)^{-1} X^\top Y.$$

When there are fewer equations than unknowns ( $n < k$ ), the system is **underdetermined**. In this case:

- There are **infinitely many exact solutions**.
- $X^\top X$  is singular and not invertible.
- The normal equation formula does not apply.

The **Moore–Penrose pseudoinverse**  $X^+$  provides a unified solution for **any matrix**. It allows us to write:

$$W = X^+ Y$$

as a generalized inverse-based solution. Special cases:

- If  $n > k$  and  $X$  has full column rank:

$$X^+ = (X^\top X)^{-1} X^\top$$

- If  $n < k$  and  $X$  has full row rank:

$$X^+ = X^\top (XX^\top)^{-1}$$

For underdetermined systems, the full solution set can be decomposed as:

$$W = W_p + W_h$$

where:

- $W_p = X^\top (XX^\top)^{-1}Y$  is a **particular solution**
- $W_h \in \text{Null}(X)$  is any vector in the null space of  $X$

More explicitly:

$$W = X^\top (XX^\top)^{-1}Y - (I - X^\top (XX^\top)^{-1}X)z, \quad z \in \mathbb{R}^k$$

Among all possible solutions, the pseudoinverse selects a very special one:

| The solution with **minimum Euclidean norm**.

That is:

$$W^* = \arg \min_W \|W\|_2 \quad \text{s.t. } XW = Y.$$

This solution is exactly:

$$W^* = X^+Y.$$

All other solutions differ by adding null-space components that increase the norm.



**Question 1** Assume we have  $N$  training samples  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$  where for each sample  $i \in \{1, \dots, N\}$  we have that  $\mathbf{x}_i \in \mathbb{R}^d$  and  $y_i \in \{-1, 1\}$ . We want to classify the dataset using the exponential loss  $L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \exp(-y_i \mathbf{x}_i^\top \mathbf{w})$  for  $\mathbf{w} \in \mathbb{R}^d$ . Which of the following statements is **true**:

- The loss function  $L$  is non-convex in  $\mathbf{w}$ .
- There exists a vector  $\mathbf{w}^*$  such that  $L(\mathbf{w}^*) = 0$ .
- If I find a vector  $\mathbf{w}^*$  such that  $L(\mathbf{w}^*) < 1/N$ , then  $\mathbf{w}^*$  linearly separates my dataset.
- None of the statements are true.
- This corresponds to doing logistic regression as seen in class.

**Solution:**  $L(w^*) < 1/N$  implies  $\exp(-y_i \mathbf{x}_i^\top w^*) < 1 \forall i$ , which means that  $y_i \mathbf{x}_i^\top w^* > 0 \forall i$ .

即预测值的符号与真实标签一致，所有样本都被正确分类。因此数据是**线性可分的**。



**Question 2** Assume we have  $N$  training samples  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$  where for each sample  $i \in \{1, \dots, N\}$  we have that  $\mathbf{x}_i \in \mathbb{R}^d$  and  $y_i \in \mathbb{R}$ . For  $\lambda \geq 0$ , we consider the following loss:

$$L_\lambda(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{x}_i^\top \mathbf{w})^2 + \lambda \|\mathbf{w}\|_2,$$

and let  $C_\lambda = \min_{\mathbf{w} \in \mathbb{R}^d} L_\lambda(\mathbf{w})$  denote the optimal loss value. Which of the following statements is **true**:

- $C_\lambda$  is a non-increasing function of  $\lambda$ .
- For  $\lambda = 0$ , the loss  $L_0$  is convex and has a unique minimizer.
- None of the statements are true.
- $C_\lambda$  is a non-decreasing function of  $\lambda$ .

**Solution:** For  $\lambda_1 < \lambda_2$ ,  $L_{\lambda_1}(\mathbf{w}) \leq L_{\lambda_2}(\mathbf{w}) \forall \mathbf{w}$ , which means that  $C_{\lambda_1} \leq C_{\lambda_2}$ .



**Question 18** Let  $\mathbf{x}, \mathbf{w}, \boldsymbol{\delta} \in \mathbb{R}^d$ ,  $y \in \{-1, 1\}$ , and  $\varepsilon \in \mathbb{R}_{>0}$  be an arbitrary positive value. Which of the following is NOT true in general:

- $\operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \log_2(1 + \exp(-y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}))) = \operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \mathbf{1}_{y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}) \leq 0}$
- $\operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \log_2(1 + \exp(-y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}))) = \operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} 1 - \tanh(y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}))$   
where  $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$
- $\operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \log_2(1 + \exp(-y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}))) = \operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \exp(-y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}))$
- $\operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \log_2(1 + \exp(-y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta}))) = \operatorname{argmin}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} y\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta})$

**Solution:**  $\operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \log_2(1 + \exp(-\langle \mathbf{w}, \mathbf{x} + \boldsymbol{\delta} \rangle)) = \operatorname{argmax}_{\|\boldsymbol{\delta}\|_2 \leq \varepsilon} \mathbf{1}_{\langle \mathbf{w}, \mathbf{x} + \boldsymbol{\delta} \rangle \leq 0}$  is not true in general since the right-hand side can have multiple maximizers. To correct it, it suffices to substitute  $=$  with  $\in$ . As for the other choices, all the right-hand sides contain losses which are monotonically decreasing in the margin  $\mathbf{w}^\top(\mathbf{x} + \boldsymbol{\delta})$  and thus lead to the same maximizer.

## Optimization

**Bias–variance tradeoff:** Increasing model complexity reduces bias but may increase variance.

What is the primary purpose of introducing nonlinear feature mappings  $\phi(x)$ ?  
To allow nonlinear decision boundaries while keeping optimization linear in parameters

In higher dimensions, **the direction of steepest descent is given by the negative gradient**.

In a contour plot of a loss function, each **contour line** represents points with identical loss values.

Geometrically, the steepest descent direction is always **orthogonal** to the contour lines of the loss function.

**Directional derivatives.** Let  $f(x, y)$  be a scalar-valued function and let  $\mathbf{u}$  be a unit vector in the plane. We examine how the function value changes if we move a small distance in direction  $\mathbf{u}$ . Using a first-order Taylor approximation:

$$\Delta z \approx \nabla f(a, b) \cdot \mathbf{u}, \Delta t$$

This shows that **the rate of change of the function** along direction  $\mathbf{u}$  depends on the dot product between the gradient and the direction vector.

Formally, the directional derivative of  $f$  at point  $\mathbf{x}$  in direction  $\mathbf{u}$  is defined as:

$$D_{\mathbf{u}}f(\mathbf{x}) = \nabla f(\mathbf{x}) \cdot \mathbf{u}$$

This quantity measures how fast the function increases or decreases when moving along  $\mathbf{u}$ . If the dot product is positive, the function increases. If it is negative, the function decreases.

Note: The requirement that  $\mathbf{u}$  be a unit vector is important because it removes dependence on step size.

Since  $\mathbf{u}$  is a unit vector, the directional derivative satisfies:

$$D_{\mathbf{u}}f = \|\nabla f\| \cos(\theta)$$

where  $\theta$  is the angle between  $\nabla f$  and  $\mathbf{u}$ . As  $\mathbf{u}$  varies,  $\cos(\theta)$  ranges between  $-1$  and  $1$ . Therefore, the maximum possible directional derivative is  $\|\nabla f\|$ , achieved when  $\mathbf{u}$  points in the same direction as the gradient.

The normalized steepest descent direction is

$$\mathbf{u} = -\frac{\nabla f}{\|\nabla f\|}$$

In practice, this normalization is omitted because the learning rate absorbs the scaling.



**Question 16** In matrix factorization for recommender systems, the goal is to approximate a partially observed user-item rating matrix  $X$  by decomposing it as the product of two lower-dimensional matrices:  $W$  (item features) and  $Z$  (user features). The observed ratings  $x_{dn}$  represent the rating of the  $n^{\text{th}}$  user for the  $d^{\text{th}}$  item, with  $\Omega$  as the set of indices for observed ratings. Which of the following is true regarding the regularized matrix factorization objective, given by:

$$\min_{W,Z} \frac{1}{2} \sum_{(d,n) \in \Omega} (x_{dn} - (WZ^\top)_{dn})^2 + \frac{\lambda_W}{2} \|W\|_F^2 + \frac{\lambda_Z}{2} \|Z\|_F^2$$

- The optimization requires filling in all missing entries in  $X$  before training.
- The Frobenius norm regularization ensures the sparsity (the number of non-zero entries) of  $W$  and  $Z$ .
- The cost function is jointly convex in  $W$  and  $Z$ .
- Regularization terms  $\frac{\lambda_W}{2} \|W\|_F^2$  and  $\frac{\lambda_Z}{2} \|Z\|_F^2$  are used to prevent overfitting to the observed ratings.

这是一个**推荐系统里的矩阵分解问题**。有一个用户-物品评分矩阵  $X$ , 但只有**一部分评分是已知的**,  $\Omega$  表示“哪些位置的评分是观测到的”, 想用两个低维矩阵去近似它: 物品特征矩阵  $W$  和用户特征矩阵  $Z$ , 用  $WZ^\top$  去拟合  $X$

**选项 1:** ~~错误~~ 错误 损失函数只对  $\Omega$  中的观测项求和。没出现的评分, 根本不进目标函数。优化时不需要、也不应该先填补缺失值。

**选项 2:** ~~错误~~ 错误 Frobenius 范数的作用是: 让参数变小但不鼓励变成 0。

**选项 3:** ~~错误~~ 错误 固定  $Z$ , 关于  $W$  是凸的, 固定  $W$ , 关于  $Z$  是凸的, 同时对  $W, Z$  优化时不是凸的。



**Question 4** Consider the loss function  $L : \mathbb{R}^d \rightarrow \mathbb{R}$ ,  $L(\mathbf{w}) = \frac{\beta}{2} \|\mathbf{w}\|^2$ , where  $\beta > 0$  is a constant. We run gradient descent on  $L$  with a stepsize  $\gamma > 0$  starting from some  $\mathbf{w}_0 \neq 0$ . Which of the statements below is true?

- Gradient descent converges in two steps for  $\gamma = \frac{1}{\beta}$  (i.e.,  $\mathbf{w}_2$  is the **first** iterate attaining the global minimum of  $L$ ).
- Gradient descent with stepsize  $\gamma = \frac{2}{\beta}$  produces iterates that diverge to infinity ( $\|\mathbf{w}_t\| \rightarrow \infty$  as  $t \rightarrow \infty$ ).
- Gradient descent converges to the global minimum for any stepsize  $\gamma > 0$ .
- Gradient descent converges to the global minimum for any stepsize in the interval  $\gamma \in (0, \frac{2}{\beta})$ .

**Solution:** The update rule is  $\mathbf{w}_{t+1} = \mathbf{w}_t - \gamma \beta \mathbf{w}_t = (1 - \gamma \beta) \mathbf{w}_t$ . Therefore we have that the sequence  $\{\|\mathbf{w}_t\|\}_t$  is given by  $\|\mathbf{w}_{t+1}\| = |1 - \gamma \beta| \|\mathbf{w}_t\| = |1 - \gamma \beta|^t \|\mathbf{w}_0\|$ . We can see that for  $\gamma = \frac{2}{\beta}$  the elements of the aforementioned sequence never move from  $\|\mathbf{w}_0\|$  (so the algorithm does not diverge to infinity for this stepsize). For  $\gamma = \frac{1}{\beta}$  the algorithm converges in one step, not two. And finally, for any  $\gamma \in (0, \frac{2}{\beta})$  the algorithm will converge to the global minimum since  $|1 - \gamma \beta| \in (0, 1)$ .

即必须满足  $|1 - \gamma \beta| < 1$ 。

## Convergence

We consider linear regression with mean squared error loss:

$$J(w) = \frac{1}{2m}(Xw - y)^T(Xw - y)$$

This objective is a **convex quadratic function**. The gradient is:

$$\nabla J(w) = \frac{1}{m}X^T(Xw - y)$$

Gradient descent updates the parameters as:

$$w_{k+1} = w_k - \alpha \nabla J(w_k)$$

The global minimizer  $w^*$  satisfies:

$$\nabla J(w^*) = 0 \quad \Rightarrow \quad w^* = (X^T X)^{-1} X^T y$$

Define the error vector:

$$e_k = w_k - w^*$$

Substituting(代入) the update rule and using  $X^T y = X^T X w^*$ , we obtain:

$$e_{k+1} = \left(I - \frac{\alpha}{m} X^T X\right) e_k$$

This shows that gradient descent behaves like a **linear dynamical system** in the error space. Convergence depends entirely on the **spectral properties** of the matrix:

$$I - \frac{\alpha}{m} X^T X$$

Using the **spectral norm**:

$$\|A\|_2 = \sup_{z \neq 0} \frac{\|Az\|_2}{\|z\|_2}$$

which is equal to **the largest singular value of A**. We obtain the bound:

$$\|e_{k+1}\|_2 \leq \left\| I - \frac{\alpha}{m} X^T X \right\|_2 \leq \|e_k\|_2$$

For convergence, we require:

$$|1 - \alpha\lambda_i| < 1 \quad \text{for all eigenvalues } \lambda_i \text{ of } \frac{1}{m}X^T X$$

which ensures **contraction of the error mapping**. This yields the fundamental step-size condition:

$$0 < \alpha < \frac{2}{\lambda_{\max}}$$

This is not a heuristic—it is a **necessary and sufficient condition** for convergence in this quadratic case.

Poor conditioning of  $X^T X$  causes slow convergence because level sets become highly elongated.

In ill-conditioned problems, tuning the learning rate alone often fails because different directions require incompatible step sizes.

Given a descent direction  $-\nabla f(w_t)$ , the **Armijo–Goldstein condition** requires:

$$J(w_t - \alpha\nabla f(w_t)) \leq J(w_t) - \frac{1}{2}\alpha\|\nabla f(w_t)\|^2$$

Interpretation: The actual decrease in loss must be at least a fixed fraction of the predicted first-order decrease. The condition guarantees sufficient decrease relative to first-order prediction.

In practice: Start with a candidate  $\alpha$ . **If the condition fails, reduce  $\alpha$**  (often by halving). Repeat until the condition is satisfied. This guarantees stability without knowing the Hessian.

Gradient descent uses only first-order information. **Newton's method** exploits **second-order structure**. Using a second-order Taylor expansion around  $x_t$ :

$$Q_t(x) = f(x_t) + \nabla f(x_t)^T(x - x_t) + \frac{1}{2}(x - x_t)^T H(x_t)(x - x_t)$$

Here,  $H(x_t)$  is the Hessian matrix. The next iterate is defined as:

$$x_{t+1} = \arg \min_x Q_t(x) \quad \Rightarrow \quad x_{t+1} = x_t - [H(x_t)]^{-1}\nabla f(x_t)$$

This step jumps directly to the minimum of the local quadratic approximation. The Newton update can fail in practice primarily because the Hessian may be ill-conditioned or indefinite.

Newton's method requires strong assumptions:

- The function must be continuous  $C^0$ ,
- Differentiable  $C^1$ ,
- And twice differentiable  $C^2$ .

Not all continuous functions are differentiable. **All differentiable functions are continuous.** These distinctions matter for optimization guarantees.

**Gradient Descent.** In stochastic gradient descent (SGD), the gradient used for each update is best described as noisy but unbiased.

## Convexity

A set  $X \subseteq \mathbb{R}^n$  is **convex** if for any two points  $x_1, x_2 \in X$ , the entire line segment between them stays inside the set:

$$\forall \lambda \in [0, 1], \quad \lambda x_1 + (1 - \lambda)x_2 \in X.$$

A function  $f : X \rightarrow \mathbb{R}$  (with **convex domain**  $X$ ) is **convex** if the function value at any point on the segment between  $x_1$  and  $x_2$  is no larger than the linear interpolation of the endpoint values:

$$\forall x_1, x_2 \in X, \forall \lambda \in [0, 1], \quad f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2).$$

Note: The domain  $X$  must be convex in the definition of a convex function, otherwise the interpolation point  $\lambda x_1 + (1 - \lambda)x_2$  may fall outside the domain.

Minimizing a **convex** function or maximizing a **concave** function guarantees that every local optimum is also a **global optimum**.

### (A) 0th-order condition (Definition / Jensen form)

$$f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2).$$

This is the core definition and does **not require derivatives**.

### (B) 1st-order condition (Supporting hyperplane)

If  $f$  is **differentiable**, convexity is equivalent to:

—

$$f(x_2) \geq f(x_1) + \nabla f(x_1)^\top (x_2 - x_1), \quad \forall x_1, x_2 \in X.$$

Meaning: the tangent plane(切平面) at  $x_1$  is a **global underestimator** of  $f$ .

### (C) 2nd-order condition (Hessian test)

If  $f$  is twice differentiable, convexity is equivalent to:

$$\nabla^2 f(x) \succeq 0 \quad \text{for all } x \in X,$$

i.e., the Hessian is **positive semidefinite(PSD)** everywhere.

If a twice-differentiable function satisfies  $\nabla^2 f(x) \succeq 0$  for all  $x \in X$ , then  $f$  is convex.

**A twice-differentiable function with a positive semidefinite Hessian everywhere is convex**, but the converse is false because convex functions need not be twice differentiable, and without a Hessian, the PSD condition is undefined.

Let  $f(x) = x^2$  with domain  $D = [-2, -1] \cup [1, 2]$ . Under standard convex analysis conventions,  $f$  is convex on each interval but not convex on  $D$  as a whole.

If  $f$  and  $g$  are convex on a convex set  $X$ , then for any  $\alpha, \beta \geq 0$ :

$$h(x) = \alpha f(x) + \beta g(x)$$

is convex on  $X$ .

If  $f$  and  $g$  are convex on  $X$ , then

$$h(x) = \max\{f(x), g(x)\}$$

is convex on  $X$ .

Let  $f \circ g = f(g(x))$ .

- If  $g$  is **affine** and  $f$  is convex, then  $f \circ g$  is convex ( $g(x) = Ax + b$ ).
- More generally, if  $f$  is **convex and non-decreasing** and  $g$  is convex, then  $f \circ g$  is convex.

This monotonicity condition prevents the composition from flipping curvature.

**If a convex function has a local minimizer, then every local minimizer is also a global minimizer.**

In general:

Global minima  $\subseteq$  Local minima  $\subseteq$  Stationary points (for smooth  $f$ ).

**Convexity collapses the hierarchy (for differentiable convex  $f$ ):**

$$\nabla f(x^*) = 0 \iff x^* \text{ is a global minimizer.}$$

A **saddle point** is best described as a stationary point that is neither a local minimum nor a local maximum.

## Convergence of Gradient Descent

### Problem Setup

We want to solve the unconstrained optimization problem

$$\min_{x \in \mathbb{R}^n} f(x)$$

using Gradient Descent (GD):

$$x_{k+1} = x_k - \alpha \nabla f(x_k),$$

where  $\alpha > 0$  is the step size and  $x^*$  denotes a global minimizer. Define the **function error**

$$e_k := f(x_k) - f(x^*).$$

### Core Assumptions

#### Convexity

A differentiable function  $f$  is convex if for all  $x, y$ :

$$f(y) \geq f(x) + \nabla f(x)^\top (y - x).$$

This implies (for minimization) that local minima are global minima, and the tangent plane(切平面) is a global under-estimator(全局下估计).

#### L-Smoothness (Lipschitz Gradient)

$f$  is **L-smooth** if its gradient is Lipschitz:

$$||\nabla f(x) - \nabla f(y)|| \leq L||x - y||.$$

A key equivalent inequality(等价不等式) (often called the **Descent Lemma(下降引理)**) is:

$$f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} ||y - x||^2.$$

**L-smoothness:**

$$f(y) - f(x) = \int_0^1 \nabla f(z(t))^T (y - x) dt \quad z(1) = y, z(0) = 0$$

$$f(y) - f(x) = \int_0^1 (\nabla f(z(t)) - \nabla f(x) + \nabla f(x))^T (y - x) dt$$

$$f(y) - f(x) = \underbrace{\int_0^1 \nabla f(x)^T (y - x) dt}_{\text{Integral A}} + \underbrace{\int_0^1 (\nabla f(z(t)) - \nabla f(x))^T (y - x) dt}_{\text{Integral B}}$$

$$f(y) - f(x) = \nabla f(x)^T (y - x) + \text{Integral B}$$

**L-smoothness:**

$$f(y) - f(x) = \int_0^1 \nabla f(z(t))^T (y - x) dt \quad z(1) = y, z(0) = 0$$

$$f(y) - f(x) = \int_0^1 (\nabla f(z(t)) - \nabla f(x) + \nabla f(x))^T (y - x) dt$$

$$f(y) - f(x) = \underbrace{\int_0^1 \nabla f(x)^T (y - x) dt}_{\text{Integral A}} + \underbrace{\int_0^1 (\nabla f(z(t)) - \nabla f(x))^T (y - x) dt}_{\text{Integral B}}$$

$$f(y) - f(x) = \nabla f(x)^T (y - x) + \text{Integral B}$$

**Cauchy-Schwarz inequality:**

$$a \cdot b \leq |a||b|$$

$$(\nabla f(z(t)) - \nabla f(x))^T (y - x) \leq (L \cdot t ||y - x||) \cdot ||y - x|| = Lt ||y - x||^2$$

$$\text{Integral B} = \int_0^1 (\nabla f(z(t)) - \nabla f(x))^T (y - x) dt \leq \int_0^1 Lt ||y - x||^2 dt$$

$$\text{Integral B} \leq \frac{L}{2} ||y - x||^2$$

## Strong Convexity

**1<sup>st</sup> order condition**

$$\underbrace{f(x) + \nabla f(x)^T(y - x) + \frac{\mu}{2} \|y - x\|^2}_{\text{Lower Bound (Floor)}} \leq f(y) \leq \underbrace{f(x) + \nabla f(x)^T(y - x) + \frac{L}{2} \|y - x\|^2}_{\text{Upper Bound (Ceiling)}}$$

**2<sup>nd</sup> order condition**  $\nabla^2 f(x) \geq \mu I$ , positive curvature at every x

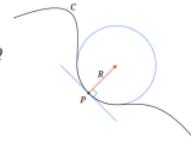
### Osculating circle [edit]

Historically, the curvature of a differentiable curve was defined through the **osculating circle**, which is the circle that best approximates the curve at a point. More precisely, given a point  $P$  on a curve, every other point  $Q$  of the curve defines a circle (or sometimes a line) passing through  $Q$  and tangent to the curve at  $P$ . The osculating circle is the **limit**, if it exists, of this circle when  $Q$  tends to  $P$ . Then the **center** and the **radius of curvature** of the curve at  $P$  are the center and the radius of the osculating circle. The curvature is the **reciprocal** of radius of curvature. That is, the curvature is

$$\kappa = \frac{1}{R},$$

where  $R$  is the radius of curvature<sup>[5]</sup> (the whole circle has this curvature, it can be read as turn  $2\pi$  over the length  $2\pi R$ ).

This definition is difficult to manipulate and to express in formulas. Therefore, other equivalent definitions have been introduced.



## Strong Convexity

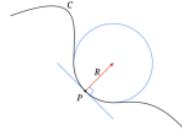
### Osculating circle [edit]

Historically, the curvature of a differentiable curve was defined through the **osculating circle**, which is the circle that best approximates the curve at a point. More precisely, given a point  $P$  on a curve, every other point  $Q$  of the curve defines a circle (or sometimes a line) passing through  $Q$  and tangent to the curve at  $P$ . The osculating circle is the **limit**, if it exists, of this circle when  $Q$  tends to  $P$ . Then the **center** and the **radius of curvature** of the curve at  $P$  are the center and the radius of the osculating circle. The curvature is the **reciprocal** of radius of curvature. That is, the curvature is

$$\kappa = \frac{1}{R},$$

where  $R$  is the radius of curvature<sup>[5]</sup> (the whole circle has this curvature, it can be read as turn  $2\pi$  over the length  $2\pi R$ ).

This definition is difficult to manipulate and to express in formulas. Therefore, other equivalent definitions have been introduced.

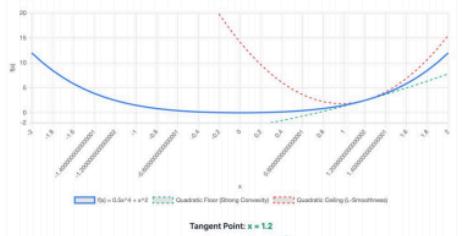


$$+ \frac{L}{2} \|y - x\|^2$$

### Example 1: A Strongly Convex & L-Smooth Function

Let's examine  $f(x) = 0.5x^2 + x^2$ . This function is strongly convex. Its curvature (second derivative) is  $f''(x) = 6x^2 + 2$ . The minimum value of the curvature is 2 (at  $x=0$ ), so it's strongly convex with  $\mu = 2$ . It is also L-smooth on the plotted interval.

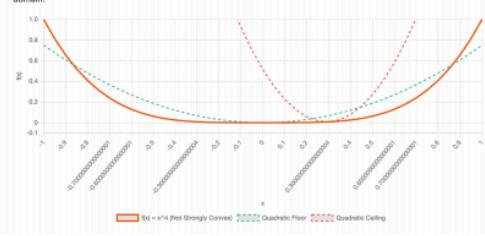
This means it's "sandwiched" between a quadratic floor (strong convexity) and a quadratic ceiling (L-smoothness). Use the slider to see how these bounds are tangent at point  $x$  but always contain the function.



### Example 2: A Convex (but not Strongly) Function

Now let's look at  $f(x) = x^2$ . This function is convex, but its curvature at the minimum,  $f''(x) = 12x^2$ , is zero, making it "too flat" to be strongly convex. The quadratic floor (green) "pokes through" the function (orange), violating the condition. The function is also not globally L-smooth, as its curvature is unbounded. However, on the interval we are plotting, it is L-smooth, allowing us to visualize the ceiling.

Use the sliders to see how no single  $\mu > 0$  can create a floor, while the ceiling always holds true on this limited domain.



## Guaranteed One-Step Progress (Choosing a Step Size)

Apply the Descent Lemma with  $x = x_k$ ,  $y = x_{k+1} = x_k - \alpha \nabla f(x_k)$ :

$$f(x_{k+1}) \leq f(x_k) + \nabla f(x_k)^T(x_{k+1} - x_k) + \frac{L}{2} \|x_{k+1} - x_k\|^2.$$

Substitute  $x_{k+1} - x_k = -\alpha \nabla f(x_k)$ :

$$\begin{aligned} f(x_{k+1}) &\leq f(x_k) - \alpha \|\nabla f(x_k)\|^2 + \frac{L\alpha^2}{2} \|\nabla f(x_k)\|^2 \\ &= f(x_k) - \alpha \left(1 - \frac{L\alpha}{2}\right) \|\nabla f(x_k)\|^2. \end{aligned}$$

So if  $\alpha \leq \frac{1}{L}$ , then  $1 - \frac{L\alpha}{2} \geq \frac{1}{2}$  and we get the clean bound (taking  $\alpha = \frac{1}{L}$ ):

$$f(x_{k+1}) \leq f(x_k) - \frac{1}{2L} \|\nabla f(x_k)\|^2.$$

This shows the objective is **monotonically decreasing** with that step size. Also rearrange it into:

$$\|\nabla f(x_k)\|^2 \leq 2L(f(x_k) - f(x_{k+1})) = 2L(e_k - e_{k+1}). \quad (1)$$

## Connecting Error to Distance (The Key Bridge)

From convexity, set  $x = x_k$ ,  $y = x^*$ :

$$\begin{aligned} f(x^*) &\geq f(x_k) + \nabla f(x_k)^\top (x^* - x_k) \\ \Rightarrow e_k &= f(x_k) - f(x^*) \leq \nabla f(x_k)^\top (x_k - x^*). \end{aligned}$$

Now expand the squared distance after one GD step:

$$\begin{aligned} \|x_{k+1} - x^*\|^2 &= \|x_k - \alpha \nabla f(x_k) - x^*\|^2 \\ &= \|x_k - x^*\|^2 - 2\alpha \nabla f(x_k)^\top (x_k - x^*) + \alpha^2 \|\nabla f(x_k)\|^2. \end{aligned}$$

Rearrange to isolate the inner product:

$$\nabla f(x_k)^\top (x_k - x^*) = \frac{1}{2\alpha} \left( \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2 \right) + \frac{\alpha}{2} \|\nabla f(x_k)\|^2.$$

Combine with  $e_k \leq \nabla f(x_k)^\top (x_k - x^*)$ :

$$e_k \leq \frac{1}{2\alpha} \left( \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2 \right) + \frac{\alpha}{2} \|\nabla f(x_k)\|^2.$$

When  $\alpha = 1/L$ :

$$e_k \leq \frac{L}{2} \left( \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2 \right) + \frac{1}{2L} \|\nabla f(x_k)\|^2. \quad (2)$$

## The Cancellation Trick (Telescoping)

Substitute equation (1) into equation (2):

$$e_k \leq \frac{L}{2} \left( \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2 \right) + (e_k - e_{k+1}).$$

Cancel  $e_k$  from both sides:

$$e_{k+1} \leq \frac{L}{2} \left( \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2 \right).$$

Sum this inequality for  $k = 0$  to  $K - 1$  (telescoping sum):

$$\begin{aligned} \sum_{k=0}^{K-1} e_{k+1} &= \sum_{k=1}^K e_k \leq \sum_{k=0}^{K-1} \frac{L}{2} \left( \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2 \right) \\ &= \frac{L}{2} \left( \|x_0 - x^*\|^2 - \|x_K - x^*\|^2 \right) \leq \frac{L}{2} \|x_0 - x^*\|^2. \end{aligned}$$

Since  $f(x_k)$  is decreasing,  $e_k$  is non-increasing, so  $e_K \leq \frac{1}{K} \sum_{k=1}^K e_k$ . Therefore:

$$e_K = f(x_K) - f(x^*) \leq \frac{L \|x_0 - x^*\|^2}{2K}.$$

To ensure  $e_K \leq \varepsilon$ , it suffices that

$$K \geq \frac{L \|x_0 - x^*\|^2}{2\varepsilon},$$

which is the classical  $O(1/\varepsilon)$  convergence rate for GD on convex, L-smooth functions.

## Strong Convexity

If  $f$  is  **$\mu$ -strongly convex**, then:

$$f(y) \geq f(x) + \nabla f(x)^\top (y - x) + \frac{\mu}{2} \|y - x\|^2,$$

equivalently (when twice differentiable)  $\nabla^2 f(x) \succeq \mu I$ . With both L-smoothness and strong convexity, GD can achieve **linear convergence** (geometric decay) rather than  $O(1/K)$ .

# Mini-batch SGD, Momentum, NAG, RMSProp, Adam

## Mini-batch SGD (The Practical Default)

We often minimize an empirical risk

$$L(w) = \frac{1}{N} \sum_{i=1}^N \ell_i(w).$$

Full-batch GD uses  $\nabla L(w)$ , which is expensive when  $N$  is large. **Mini-batch SGD** approximates the gradient using a small batch  $B$ :

$$g_t \approx \nabla L(w_t) = \frac{1}{|B|} \sum_{i \in B} \nabla \ell_i(w_t), \quad w_{t+1} = w_t - \alpha g_t.$$

## Why Stochasticity(随机性) Can Be Preferable

In convex problems, GD has clean guarantees and the landscape is “one bowl.”

In non-convex problems (typical deep nets), the landscape contains many critical points, including **saddle points**. The noise in SGD can help exploration:

- It can “kick” the iterate out of flat/unstable saddle regions.
- It can help avoid getting stuck in certain sharp basins(盆地) (informal but often observed).

## Dampen the Oscillation(抑制振荡): Momentum Methods

Plain SGD can zig-zag, especially in “ravines(深谷)” where curvature is steep in one direction and flat in another. Momentum adds a memory of past gradients to smooth updates. One common form (velocity formulation):

$$v_{t+1} = \beta v_t + g_t, \quad w_{t+1} = w_t - \alpha v_{t+1},$$

with  $v_0 = 0$  and  $\beta \in [0, 1]$  (often around 0.9–0.99).

**Interpretation:**  $v_t$  becomes an exponentially weighted moving average of gradients:

$$v_t = g_t + \beta g_{t-1} + \beta^2 g_{t-2} + \dots$$

So directions that persist get amplified; directions that alternate(反复交替) (oscillations) cancel out(相互抵消).

## Nesterov Accelerated Gradient (NAG): Look Ahead, Then Correct

Momentum can overshoot(超调) because it commits(笃信) strongly to the velocity. NAG modifies where you measure the gradient: evaluate it at a "lookahead" point. A typical NAG-like update:

$$\tilde{w}_t = w_t - \alpha \beta v_t, \quad v_{t+1} = \beta v_t + \nabla L(\tilde{w}_t), \quad w_{t+1} = w_t - \alpha v_{t+1}.$$

**Intuition:** "move in the momentum direction first, then compute the gradient there and correct." This often stabilizes training and can speed convergence compared to vanilla momentum.

## Speeding Up Convergence: Step-size Heuristics

**Move slowly in steep directions, fast in flat directions.** This motivates per-parameter adaptive learning rates: scale(调整) updates differently in each coordinate(坐标分量) depending on gradient history.

## RMSProp: Adaptive Step Sizes via Running Squared Gradients(平方梯度滑动量)

RMSProp keeps a moving average of squared gradients:

$$s_t = \beta s_{t-1} + (1 - \beta) g_t^2 \quad (\text{elementwise square}),$$

and updates

$$w_{t+1} = w_t - \alpha \frac{g_t}{\sqrt{s_t + \varepsilon}}.$$

**Effect:** if a coordinate has consistently large gradients,  $s_t$  grows and the effective step size shrinks for that coordinate; if gradients are small, steps become larger.

## Adam: Momentum + RMSProp (Plus Bias Correction)

Adam combines:

- a first-moment(一阶矩) estimate (momentum-like):

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t,$$

- a second-moment estimate (RMSProp-like):

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2.$$

Because  $m_t, v_t$  start at zero, they are biased early. Adam applies bias correction:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

Update:

$$w_{t+1} = w_t - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}.$$

Common defaults:  $\beta_1 = 0.9, \beta_2 = 0.999, \alpha$  tunable.

## Comparison at a Glance (What problem each method targets)

Momentum	RMSprop	Adam
<b>Core Idea</b> Accelerates gradient descent by adding a fraction of the past update vector to the current one.	<b>Core Idea</b> Adapts the learning rate for each parameter based on the magnitude of recent gradients.	<b>Core Idea</b> Combines the ideas of both Momentum and RMSprop.
<b>Learning Rate</b> Global (one for all parameters).	<b>Learning Rate</b> Per-parameter and adaptive.	<b>Learning Rate</b> Per-parameter and adaptive.
<b>Key Mechanism</b> Stores a moving average of gradients (first moment).	<b>Key Mechanism</b> Stores a moving average of "squared" gradients (second moment).	<b>Key Mechanism</b> Stores moving averages of both gradients and squared gradients.
<b>Best For</b> Improving convergence speed and overcoming local minima in many scenarios.	<b>Best For</b> Problems with noisy or sparse gradients, and non-stationary objectives (e.g., RNNs).	<b>Best For</b> A robust, all-purpose default optimizer that works well for a wide variety of problems.
<b>Main Weakness</b> Can struggle if parameters need vastly different learning rates.	<b>Main Weakness</b> Doesn't benefit from the acceleration that momentum provides.	<b>Main Weakness</b> Can sometimes converge to a less optimal solution than well-tuned SGD. Requires more memory.

## Takeaway Messages

**Reduced oscillation:** Momentum makes the trajectory smoother, especially in "ravines(深谷)" where one direction is steep and the other is flat. It damps(抑制) the classic SGD zig-zag.

**Faster convergence (often):** smoothing usually creates a more direct path and builds speed in consistent downhill directions.

**Nesterov's advantage (NAG):** by "looking ahead," NAG tends to brake more appropriately near the minimum, reducing overshoot and making training more stable.

**New hyperparameter:** Momentum introduces  $\beta$  (momentum coefficient(系数)). Default values like 0.9 often work well, but the *interaction* between  $\alpha$  and  $\beta$  matters—too large a learning rate plus high momentum can turn your optimizer into a pinball(弹珠) machine.

## Practical Advice

**Start with Adam** for most deep learning tasks: it's a strong baseline because it handles scaling differences and gradient noise with minimal tuning.

**Consider RMSProp** in settings that are highly non-stationary or historically tricky (e.g., some RNN / RL-style training loops). (Modern practice often still uses Adam variants here, but RMSProp remains conceptually relevant.)

**Don't forget SGD + Momentum:** when you have time to tune and you care about best final generalization, SGD+Momentum can sometimes beat Adam—especially in supervised vision-style regimes(系统) with good schedules.



**Question 3** Which of the following statements is **always** true for a real-valued data matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$  and  $\mathbf{y} \in \mathbb{R}^n$  with  $d > n$  where  $n$  is the number of data points and  $d$  represent the feature dimension?

- The sample points are linearly separable.
- The data points lie in an at most  $n$ -dimensional subspace of  $\mathbb{R}^d$ , so there is at least one non-zero direction in  $\mathbb{R}^d$  that is orthogonal to all these points.
- The weights  $\mathbf{w}$  can be computed with least-squares linear regression as follows:  $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ .
- $\mathbf{X}^\top \mathbf{X}$  has exactly  $d - n$  eigenvectors with eigenvalue zero.

**Solution:** In the case where  $d > n$  (more features than samples),  $\mathbf{X}^\top \mathbf{X}$  is not invertible because it is a  $d \times d$  matrix with rank at most  $n$  since  $\mathbf{X}$  has only  $n$  rows. When  $\mathbf{X}^\top \mathbf{X}$  is not full rank, its inverse does not exist, hence the least squares solution cannot be computed with the given formula. The matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$  can have at most  $n$  linearly independent rows, so its rank is at most  $n$ . As a result, the rank of the  $d \times d$  matrix  $\mathbf{X}^\top \mathbf{X}$  is also limited to  $n$ , implying that there can be **at most**  $n$  non-zero eigenvalues. Consequently, there are **at least**  $d - n$  zero eigenvalues and the statement cannot be considered definitively true. Linear separability of the sample points is not guaranteed simply by having  $d > n$ , it depends on the specific data. This remaining option is true.

给定：特征数多于样本数，这是典型的 **高维小样本 (overparameterized)** 场景。

### ✗ 选项 1

**线性可分性不是只由维度决定的。** 反例很简单：所有点完全重合；或者标签随机但特征退化。线性可分性取决于**数据分布和标签**，不是“维度是否大于样本数”。

### ✓ 选项 2

$X$  的行向量最多  $n$  个线性无关，所以所有数据点 **最多张成一个  $n$  维子空间**。整个空间是  $\mathbb{R}^d$ ，而子空间维数  $\leq n$ ，且  $d > n \rightarrow$  **必然存在非零的正交补空间。**

存在  $v \in \mathbb{R}^d, v \neq 0$ ，使得： $Xv = 0$ 。也就是说：这个方向对所有数据点“不可见”；数据在这个方向上的投影全为 0。

### ✗ 选项 3

$$\mathbf{X}^\top \mathbf{X} \in \mathbb{R}^{d \times d}$$

但：

$$\text{rank}(\mathbf{X}^\top \mathbf{X}) = \text{rank}(\mathbf{X}) \leq n < d$$

**不是满秩矩阵，不可逆，所以这个公式根本不存在。**

正确做法应该是：Moore–Penrose 伪逆或加正则项 (Ridge)。

### ✗ 选项 4

我们只知道： $\text{rank}(\mathbf{X}^\top \mathbf{X}) \leq n$ ，因此：非零特征值 **最多  $n$  个** 零特征值 **至少  $d - n$  个**。但注意：至少 ≠ 恰好。如果： $X$  的行本身就线性相关，那

么  $\text{rank}(X) < n$ , 零特征值会 **多于  $(d-n)$** 。